

```
In [100]: import numpy as np
import sklearn
from sklearn import model_selection
import pandas as pd
import sys
import matplotlib.pyplot as plt
```

```
In [101]: disease_data= pd.read_csv(r'D:\Google_Download\Artificial Intelligence\Assignment\assignment 2-supp.csv')
```

```
In [102]: disease_data.head()
```

```
Out[102]:
```

	Pregnancies	Glucose	BloodPressure	SkinThickness	Insulin	BMI	DiabetesPedigreeFunction	Age	Outcome
0	6	148	72	35	0	33.6	0.627	50	1
1	1	85	66	29	0	26.6	0.351	31	0
2	8	183	64	0	0	23.3	0.672	32	1
3	1	89	66	23	94	28.1	0.167	21	0
4	0	137	40	35	168	43.1	2.288	33	1

```
In [103]: disease_data.shape
```

```
Out[103]: (768, 9)
```

```
In [104]: train, test = model_selection.train_test_split(disease_data, test_size=0.2, random_state=1234)
```

```
In [108]: real_train, validation=model_selection.train_test_split(train, test_size=0.25, random_state=1234)
```

```
In [105]: train.shape
```

```
Out[105]: (614, 9)
```

```
In [106]: test.shape
```

```
Out[106]: (154, 9)
```

```
In [109]: real_train.shape
```

```
Out[109]: (460, 9)
```

```
In [110]: validation.shape
```

```
Out[110]: (154, 9)
```

```
In [111]: real_train_y=real_train.Outcome
real_train_x=real_train.drop("Outcome",axis=1)
validation_y=validation.Outcome
validation_x=validation.drop("Outcome",axis=1)
test_y=test.Outcome
test_x=test.drop("Outcome",axis=1)
```

```
In [128]: class LogisticRegression:
def __init__(self, lr=0.01, num_iter=10000, fit_intercept=True, verbose=True):
    self.lr = lr
    self.num_iter = num_iter
    self.fit_intercept = fit_intercept
    self.verbose=verbose
    self.loss_value=[]

def __add_intercept(self, X):
    intercept = np.ones((X.shape[0], 1))
    # add the bias for X
    return np.concatenate((intercept, X), axis=1)

def __sigmoid(self, z):
    return 1 / (1 + np.exp(-z))

def __loss(self, h, y):
    # Loss function-> 1/m*(-y * np.log(h(theta' * x)) - (1 - y) * np.log(1 - h(theta' * x)))
    return (-y * np.log(h) - (1 - y) * np.log(1 - h)).mean()

def fit(self, X, y):
    if self.fit_intercept:
        X = self.__add_intercept(X)
    # weights initialization -> can use all 0 initialization, but here I use random initialization to eliminate the coincidence
    #self.theta = np.random.randn(X.shape[1])
    self.theta=np.zeros(X.shape[1])


    for i in range(self.num_iter):
        z = np.dot(X, self.theta)
        h = self.__sigmoid(z)
        tmp=self.__loss(h, y)
        if tmp!=float('inf'):
            self.loss_value.append(self.__loss(h, y)) # append the loss value to see whether it reached overfit
        # Batch gradient descent
        # Calculate the gradient:
        gradient = np.dot(X.T, (h - y)) / y.size
        self.theta -= self.lr * gradient
        if(self.verbose == True and i % self.num_iter == 0):
            # Print the final loss
            z = np.dot(X, self.theta)
            h = self.__sigmoid(z)
            print(f'loss: {self.__loss(h, y)} \t')

def predict_prob(self, X):
    # probability calculation function, get the final theta value, and predict the probability of each input vector
    if self.fit_intercept:
        X = self.__add_intercept(X)
    return self.__sigmoid(np.dot(X, self.theta))

def predict(self, X, threshold):
    # Predict function: if the p is larger then threshold, then predict it to be 1. Otherwise 0.
    return self.predict_prob(X) >= threshold
```



```
In [150]: plt.plot(np.linspace(0.001,0.002,10),acu_result_learning_rate)
plt.show() # ensure the iteration number = 7500
```



```
In [144]: train_y=train.Outcome
train_x=train.drop("Outcome",axis=1)
```

```
In [158]: model = LogisticRegression(lr=0.0016, num_iter=7500)
model.fit(train_x, train_y)
```

loss: 1.7242433603491103

<ipython-input-128-bc9fba4b6c8>:18: RuntimeWarning: divide by zero encountered in log
return (-y * np.log(h) - (1 - y) * np.log(1 - h)).mean()

```
In [173]: preds = model.predict(test_x, 0.95)
# accuracy
(preds == test_y).mean()
```

Out[173]: 0.6233766233766234

```
In [169]: model.theta
```

Out[169]: array([[-0.43852211, 0.73429886, 0.04561872, -0.17067624, -0.05802802,
 -0.01903002, 0.0093603 , 0.09285052, -0.09620783])

```
In [ ]:
```