

CS 188: Artificial Intelligence

Constraint Satisfaction Problems



Instructor: Anca Dragan

University of California, Berkeley


[These slides adapted from Dan Klein and Pieter Abbeel]

Constraint Satisfaction Problems

N variables

domain D

constraints



states

*partial
assignment*

goal test

*complete; satisfies
constraints*


successor

function

assign an unassigned variable

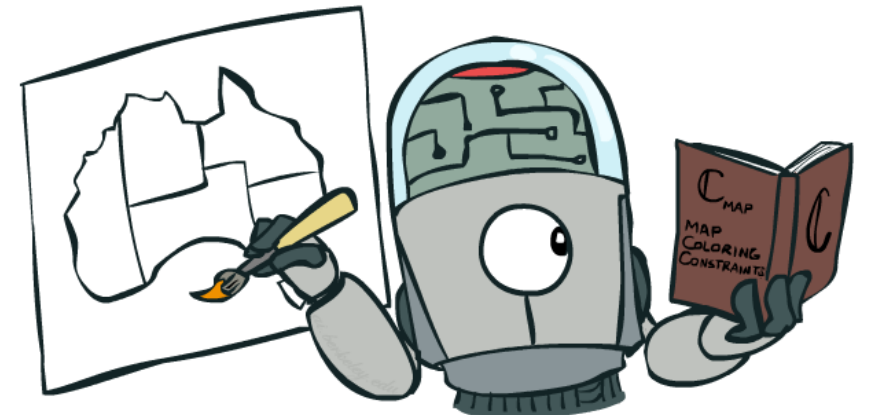

What is Search For?

- Assumptions about the world: a single agent, deterministic actions, fully observed state, discrete state space
- Planning: sequences of actions
 - **The path** to the goal is the important thing
 - Paths have various costs, depths
 - **Heuristics** give problem-specific guidance
- Identification: assignments to variables
 - **The goal itself is important, not the path**
 - **All paths at the same depth (for some formulations)**
 - CSPs are specialized for identification problems




Constraint Satisfaction Problems

- Standard search problems:
 - State is a “black box” : arbitrary data structure
 - Goal test can be any function over states
 - Successor function can also be anything
- Constraint satisfaction problems (CSPs):
 - A **special subset** of search problems
 - **State** is defined by **variables X_i** with **values** from a **domain D** (sometimes D depends on i)
 - **Goal test** is a **set of constraints** specifying allowable combinations of values for subsets of variables
- Allows useful **general-purpose algorithms** with **more power** than standard search algorithms





CSP Examples




Example: Map Coloring

- Variables WA, NT, Q, NSW, V, SA, T
- Domains: $D = \{\text{red, green, blue}\}$
- Constraints: adjacent regions must have different colors
 - Implicit: $WA \neq NT$
 - Explicit: $(WA, NT) \in \{(\text{red, green}), (\text{red, blue}), \dots\}$
- **Solutions** are *assignments satisfying all constraints*, e. g. :
 - $\{WA=\text{red}, NT=\text{green}, Q=\text{red}, NSW=\text{green}, V=\text{red}, SA=\text{blue}, T=\text{green}\}$




Constraint Graphs





Constraint Graphs

- Binary CSP: each constraint relates (at most) two variables
- Binary constraint graph: nodes are variables, arcs show constraints
- General-purpose CSP algorithms use the **graph structure** to **speed up search**. E. g., Tasmania is an independent subproblem!



Example: N-Queens

- Formulation 1:
 - Variables: X_{ij}
 - Domains: $\{0, 1\}$
 - Constraints



$$\forall i, j, k \quad (X_{ij}, X_{ik}) \in \{(0, 0), (0, 1), (1, 0)\}$$

$$\forall i, j, k \quad (X_{ij}, X_{kj}) \in \{(0, 0), (0, 1), (1, 0)\}$$

$$\forall i, j, k \quad (X_{ij}, X_{i+k, j+k}) \in \{(0, 0), (0, 1), (1, 0)\}$$

$$\forall i, j, k \quad (X_{ij}, X_{i+k, j-k}) \in \{(0, 0), (0, 1), (1, 0)\}$$

$$\sum_{i,j} X_{ij} = N$$

Example: N-Queens

- Formulation 2:

- Variables: Q_k


- Domains: $\{1, 2, 3, \dots, N\}$

- Constraints:

Implicit: $\forall i, j \text{ non-threatening}(Q_i, Q_j)$

Explicit: $(Q_1, Q_2) \in \{(1, 3), (1, 4), \dots\}$

...



Screenshot of Demo N-Queens

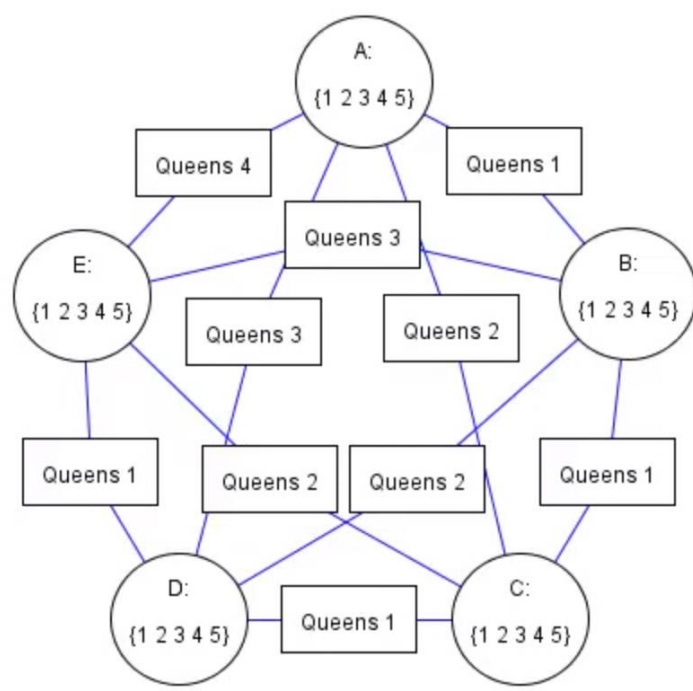
CSP Applet Version 4.6.1 --- fiveQueens.xml

File Edit View CSP Options Help

Fine Step Step Auto Arc-Consistency AutoSolve Stop Step Back Backtrack Reset

Create Solve

Click on a variable to split its domain.
Click on a constraint to reorder its variables.
Click on an arc to make it arc-consistent.



Note1 - Windows Journal

File Edit View Insert Actions Tools Help

5-QUEENS

	1	2	3	4	5
A					
B					
C					
D					
E					

1/1

Example: Cryptarithmic

- Variables:

$F T U W R O X_1 X_2 X_3$

- Domains:

$\{0, 1, 2, 3, 4, 5, 6, 7, 8, 9\}$

- Constraints:


$\text{alldiff}(F, T, U, W, R, O)$

$O + O = R + 10 \cdot X_1$


...

$$\begin{array}{r} T W O \\ + T W O \\ \hline F O U R \end{array}$$

X_1



Example: Sudoku



- Variables:
 - Each (open) square
- Domains:
 - $\{1, 2, \dots, 9\}$
- Constraints:

9-way alldiff for each column



9-way alldiff for each row

9-way alldiff for each region

(or can have a bunch of pairwise inequality constraints)

Example: The Waltz Algorithm



- The **Waltz algorithm** is for interpreting line drawings of solid polyhedra as 3D objects
- An early example of an AI computation posed as a CSP



- **Approach:**
 - Each intersection is a variable
 - Adjacent intersections impose constraints on each other
 - Solutions are physically realizable 3D interpretations


Example: The Waltz Algorithm

- The Waltz algorithm is for interpreting line drawings of solid polyhedra as 3D objects
- An early example of an AI computation posed as a CSP



- Approach:
 - Each intersection is a variable
 - Adjacent intersections impose constraints on each other
 - Solutions are physically realizable 3D interpretations

Varieties of CSPs and Constraints



Varieties of CSPs

○ Discrete Variables

○ Finite domains


- Size d means $O(d^n)$ complete assignments
- E.g., Boolean CSPs, including **Boolean satisfiability** (NP-complete)

○ Infinite domains (integers, strings, etc.)

- E.g., job scheduling, variables are start/end times for each job
- **Linear constraints solvable, nonlinear undecidable**


○ Continuous variables

- E.g., start/end times for Hubble Telescope observations
- Linear constraints **solvable in polynomial time** by LP methods (see cs170 for a bit of this theory)




Varieties of Constraints

- Varieties of Constraints
 - Unary constraints involve a **single variable** (equivalent to reducing domains), e.g.:
 $SA \neq \text{green}$
 - Binary constraints involve **pairs of variables**, e.g.:
 $SA \neq WA$
 - Higher-order constraints involve **3 or more variables**:
e.g., cryptarithmic column constraints
- **Preferences (soft constraints)**:
 - E.g., red is better than green
 - Often **representable by a cost** for each variable assignment
 - Gives constrained optimization problems
 - (We'll ignore these until we get to Bayes')



Real-World CSPs

- Assignment problems: e.g., who teaches what class
- Timetabling problems: e.g., which class is offered when and where?
- Hardware configuration
- Transportation scheduling
- Factory scheduling
- Circuit layout
- Fault diagnosis
- ... lots more!




- Many real-world problems involve real-valued variables...

Solving CSPs



Standard Search Formulation


- Standard search formulation of CSPs
- States defined by the values assigned so far (partial assignments)
 - Initial state: the empty assignment, $\{\}$
 - Successor function: assign a value to an unassigned variable
 - Goal test: the current assignment is complete and satisfies all constraints
- We'll start with the straightforward, naïve approach,



Search Methods

- What would BFS do?


$\{WA=g\}$ $\{WA=r\}$ $\{ \dots \}$ $\{NT=g\}$ \dots



Search Methods


- What would BFS do?

- What would DFS do?
 - let's see!




- What problems does naïve search have?

Video of Demo Coloring -- DFS




Backtracking Search





Backtracking Search

- Backtracking search is the **basic uninformed algorithm** for solving CSPs → 走一步看一步
- Idea 1: One variable at a time
 - Variable assignments are commutative, so fix ordering → **better branching factor!**
 - I.e., [WA = red then NT = green] same as [NT = green then WA = red]
 - Only need to consider assignments to **a single variable at each step**
- Idea 2: Check constraints as you go
 - I.e. consider only values which do not conflict previous assignments
 - Might have to **do some computation to check the constraints**
 - “Incremental goal test”
- Depth-first search with these two improvements is called *backtracking search* (not the best name)
- Can solve n-queens for $n \approx 25$




Backtracking Example



[Demo: coloring -- backtracking]

Video of Demo Coloring - Backtracking



Backtracking Search


```
function BACKTRACKING-SEARCH(csp) returns solution/failure
  return RECURSIVE-BACKTRACKING({ }, csp)

function RECURSIVE-BACKTRACKING(assignment, csp) returns soln/failure
  if assignment is complete then return assignment
  var ← SELECT-UNASSIGNED-VARIABLE(VARIABLES[csp], assignment, csp)
  for each value in ORDER-DOMAIN-VALUES(var, assignment, csp) do
    if value is consistent with assignment given CONSTRAINTS[csp] then
      add {var = value} to assignment
      result ← RECURSIVE-BACKTRACKING(assignment, csp)
      if result ≠ failure then return result
      remove {var = value} from assignment
  return failure
```

- Backtracking = DFS + **variable-ordering** + fail-on-violation
- What are the choice points?

Improving Backtracking

- General-purpose ideas give huge gains in speed
- Ordering:
 - Which variable should be assigned next?
 - In what order should its values be tried?
- Filtering: Can we detect inevitable failure early?




Filtering




Keep track of domains for unassigned variables and cross off bad options

Filtering: Forward Checking

- Filtering: Keep track of domains for unassigned variables and cross off bad options
- Forward checking: **Cross off values** that violate a constraint when added to the existing assignment




Video of Demo Coloring - Backtracking with Forward Checking



Filtering: Constraint Propagation

- Forward checking propagates information from assigned to unassigned variables, but doesn't provide early detection for all failures:






WA	NT	Q	NSW	V	SA
<div><div></div><div></div><div></div></div>	<div><div></div><div></div><div></div></div>	<div><div></div><div></div><div></div></div>	<div><div></div><div></div><div></div></div>	<div><div></div><div></div><div></div></div>	<div><div></div><div></div><div></div></div>
<div><div></div><div></div><div></div></div>	<div><div></div><div></div><div></div></div>	<div><div></div><div></div><div></div></div>	<div><div></div><div></div><div></div></div>	<div><div></div><div></div><div></div></div>	<div><div></div><div></div><div></div></div>
<div><div></div><div></div><div></div></div>	<div><div></div><div></div><div></div></div>	<div><div></div><div></div><div></div></div>	<div><div></div><div></div><div></div></div>	<div><div></div><div></div><div></div></div>	<div><div></div><div></div><div></div></div>

- NT and SA cannot both be blue!
- Why didn't we detect this yet?
- Constraint propagation*: reason from constraint to constraint

Consistency of A Single Arc

- An arc $X \rightarrow Y$ is **consistent** iff for *every* x in the tail there is *some* y in the head which could be assigned without violating a constraint





Forward checking?

Enforcing consistency of arcs pointing to each new assignment

Arc Consistency of an Entire CSP

- A simple form of propagation makes sure **all** arcs are consistent:



- Important: If X loses a value, neighbors of X need to be rechecked!
- Arc consistency **detects failure earlier** than forward checking
- Can be run as a preprocessor or after each assignment
- What's the downside of enforcing arc consistency?

*Remember: Delete
from the tail!*

Enforcing Arc Consistency in a CSP

function AC-3(*csp*) **returns** the CSP, possibly with reduced domains

inputs: *csp*, a binary CSP with variables $\{X_1, X_2, \dots, X_n\}$

local variables: *queue*, a queue of arcs, initially all the arcs in *csp*

while *queue* is not empty **do**

$(X_i, X_j) \leftarrow \text{REMOVE-FIRST}(\text{queue})$

if REMOVE-INCONSISTENT-VALUES(X_i, X_j) **then**

for each X_k **in** NEIGHBORS[X_i] **do**

 add (X_k, X_i) to *queue*

Queue 's head has at most d values, then remove-inconsistent-value need to execute $d \cdot n$ times, then Neighbors[X_i] has at most d values, then we may add $d \cdot n \cdot d$ times, for the fact that the while loop can run at most n times, thus we have $O(n^2 d^2)$

function REMOVE-INCONSISTENT-VALUES(X_i, X_j) **returns** true iff succeeds

$\text{removed} \leftarrow \text{false}$

for each x **in** DOMAIN[X_i] **do**

if no value y in DOMAIN[X_j] allows (x, y) to satisfy the constraint $X_i \leftrightarrow X_j$



then delete x from DOMAIN[X_i]; $\text{removed} \leftarrow \text{true}$

return *removed*

- Runtime: $O(n^2 d^3)$, can be reduced to $O(n^2 d^2)$
- ... but detecting all possible future problems is NP-hard - why?


Limitations of Arc Consistency

- After enforcing arc consistency:
 - Can have one solution left
 - Can have multiple solutions left
 - Can have no solutions left (and not know it)
- Arc consistency still runs inside a backtracking search!




[Demo: coloring -- forward checking]
[Demo: coloring -- arc consistency]

Video of Demo Coloring - Backtracking with Forward Checking - Complex Graph



Video of Demo Coloring - Backtracking with Arc Consistency - Complex Graph





K-Consistency

- Increasing degrees of consistency
 - 1-Consistency (Node Consistency): Each single node's domain has a value which meets that node's unary constraints
 - 2-Consistency (Arc Consistency): For each pair of nodes, any consistent assignment to one can be extended to the other
 - K-Consistency: For each k nodes, any consistent assignment to k-1 can be extended to the kth node.

- Higher k more expensive to compute

- (You need to know the k=2 case: arc consistency)



Strong K-Consistency

- Strong k-consistency: also $k-1$, $k-2$, \dots 1 consistent
- Claim: strong n-consistency means we can solve without backtracking!
- Why?
 - Choose any assignment to any variable
 - Choose a new variable
 - By 2-consistency, there is a choice consistent with the first
 - Choose a new variable
 - By 3-consistency, there is a choice consistent with the first 2
 - \dots
- Lots of middle ground between arc consistency and n-consistency! (e. g. $k=3$, called path consistency)